

Principle of Locality: Memory Hierarchies

- Text and data are not accessed randomly
- **Temporal** locality
 - Recently accessed items will be accessed in the near future (e.g., code in loops, top of stack)
- **Spatial** locality
 - Items at addresses close to the addresses of recently accessed items will be accessed in the near future (sequential code, elements of arrays)
- Leads to memory hierarchy at two main interface levels:
 - Processor - Main memory -> Introduction of **caches**
 - Main memory - Secondary memory -> **Virtual memory** (paging systems)

Processor - Main Memory Hierarchy

- **Registers**: Those visible to ISA + those renamed by hardware
- **(Hierarchy of) Caches**: plus their enhancements
 - Write buffers, victim caches etc...
- **TLB's** and their management
- Virtual memory system (O.S. level) and hardware assists (page tables)
- **Inclusion of information** (or space to gather information) level per level
 - Almost always true

Questions that Arise at Each Level

- What is the unit of information transferred from level to level ?
 - Word (byte, double word) to/from a register
 - Block (line) to/from cache
 - Page table entry + misc. bits to/from TLB
 - Page to/from disk
- When is the unit of information transferred from one level to a lower level in the hierarchy?
 - Generally, on demand (cache miss, page fault)
 - Sometimes earlier (prefetching)

Questions that Arise at Each Level (c'ed)

- Where in the hierarchy is that unit of information placed?
 - For registers, directed by ISA and/or register renaming method
 - For caches, in general in L1
 - Possibility of hinting to another level (Itanium) or of bypassing the cache entirely, or to put in special buffers
- How do we find if a unit of info is in a given level of the hierarchy?
 - Depends on mapping;
 - Use of hardware (for caches/TLB) and software structures (page tables)

Questions that Arise at Each Level (c'ed)

- What happens if there is no room for the item we bring in?
 - Replacement algorithm; depends on organization
- What happens when we change the contents of the info?
 - i.e., what happens on a write?

Caches (on-chip, off-chip)

- Caches consist of a **set of entries** where each entry has:
 - **line** (or **block**) of data: information contents
 - **tag**: allows to recognize if the block is there
 - **status bits**: valid, dirty, state for multiprocessors etc.
- Cache Geometries
 - Capacity (or size) of a cache: number of lines * line size, i.e., the cache **metadata** (tag + status bits) is not counted
 - Associativity
 - Line size

Cache Organizations

- Direct-mapped
- Set-associative
- Fully-associative

Cache Hit or Cache Miss?

- How to detect if a memory address (a byte address) has a valid image in the cache:
- Address is decomposed in 3 fields:
 - *line offset* or *displacement* (depends on line size)
 - *index* (depends on number of sets and set-associativity)
 - *tag* (the remainder of the address)
- The tag array has a width equal to *tag*

Hit Detection

tag	index	displ.
-----	-------	--------

Example: cache capacity C , line size b

Direct mapped: $\text{displ} = \log_2 b$; $\text{index} = \log_2(C/b)$; $\text{tag} = 32 - \text{index} - \text{displ}$

N -way S.A: $\text{displ} = \log_2 b$; $\text{index} = \log_2(C/bN)$; $\text{tag} = 32 - \text{index} - \text{displ}$

So what does it mean to have 3-way ($N=3$) set-associativity?

Replacement Algorithm

- None for direct-mapped
- Random or LRU or pseudo-LRU for set-associative caches
 - Not an important factor for performance for low associativity. Can become important for large associativity and large caches